

Social network Graph Link Prediction - Facebook Challenge

```
In [2]: 1 #Importing Libraries
2 # please do go through this python notebook:
3 import warnings
4 warnings.filterwarnings("ignore")
5
6 import csv
7 import pandas as pd#pandas to create small dataframes
8 import datetime #Convert to unix time
9 import time #Convert to unix time
10 # if numpy is not installed already : pip3 install numpy
11 import numpy as np#Do arithmetic operations on arrays
12 # matplotlib: used to plot graphs
13 import matplotlib
14 import matplotlib.pyplot as plt
15 import seaborn as sns#Plots
16 from matplotlib import rcParams#Size of plots
17 from sklearn.cluster import MiniBatchKMeans, KMeans#Clustering
18 import math
19 import pickle
20 import os
21 # to install xgboost: pip3 install xgboost
22 import xgboost as xgb
23
24 import warnings
25 import networkx as nx
26 import pdb
27 import pickle
28 from pandas import HDFStore, DataFrame
29 from pandas import read_hdf
30 from scipy.sparse.linalg import svds, eigs
31 import gc
32 from tqdm import tqdm
33 from sklearn.ensemble import RandomForestClassifier
34 from sklearn.metrics import f1_score
```

```
In [7]: 1 #reading
        2 from pandas import read_hdf
        3 df_final_train = read_hdf('data/fea_sample/storage_sample_stage4.h5', 'train_df',mode='r')
        4 df_final_test = read_hdf('data/fea_sample/storage_sample_stage4.h5', 'test_df',mode='r')
```

```
In [17]: 1 y_train = df_final_train.indicator_link
        2 y_test = df_final_test.indicator_link
```

```
In [ ]: 1
```

```
In [18]: 1 df_final_train.drop(['source_node', 'destination_node', 'indicator_link'],axis=1,inplace=True)
        2 df_final_test.drop(['source_node', 'destination_node', 'indicator_link'],axis=1,inplace=True)
```

```

In [0]: 1 estimators = [10,50,100,250,450]
        2 train_scores = []
        3 test_scores = []
        4 for i in estimators:
        5     clf = RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
        6         max_depth=5, max_features='auto', max_leaf_nodes=None,
        7         min_impurity_decrease=0.0, min_impurity_split=None,
        8         min_samples_leaf=52, min_samples_split=120,
        9         min_weight_fraction_leaf=0.0, n_estimators=i, n_jobs=-1, random_state=25, verbose=0, warm_start=False)
        10     clf.fit(df_final_train, y_train)
        11     train_sc = f1_score(y_train, clf.predict(df_final_train))
        12     test_sc = f1_score(y_test, clf.predict(df_final_test))
        13     test_scores.append(test_sc)
        14     train_scores.append(train_sc)
        15     print('Estimators = ', i, 'Train Score', train_sc, 'test Score', test_sc)
        16 plt.plot(estimators, train_scores, label='Train Score')
        17 plt.plot(estimators, test_scores, label='Test Score')
        18 plt.xlabel('Estimators')
        19 plt.ylabel('Score')
        20 plt.title('Estimators vs score at depth of 5')

```

```


Estimators = 10 Train Score 0.9063252121775113 test Score 0.8745605278006858
Estimators = 50 Train Score 0.9205725512208812 test Score 0.9125653355634538
Estimators = 100 Train Score 0.9238690848446947 test Score 0.9141199714153599
Estimators = 250 Train Score 0.9239789348046863 test Score 0.9188007232664732
Estimators = 450 Train Score 0.9237190618658074 test Score 0.9161507685828595

```

```

Out[6]: Text(0.5,1,'Estimators vs score at depth of 5')

```



```


In [0]: 1 depths = [3,9,11,15,20,35,50,70,130]
        2 train_scores = []
        3 test_scores = []
        4 for i in depths:
        5     clf = RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
        6         max_depth=i, max_features='auto', max_leaf_nodes=None,
        7         min_impurity_decrease=0.0, min_impurity_split=None,
        8         min_samples_leaf=52, min_samples_split=120,
        9         min_weight_fraction_leaf=0.0, n_estimators=115, n_jobs=-1, random_state=25, verbose=0, warm_start=
10     clf.fit(df_final_train,y_train)
11     train_sc = f1_score(y_train,clf.predict(df_final_train))
12     test_sc = f1_score(y_test,clf.predict(df_final_test))
13     test_scores.append(test_sc)
14     train_scores.append(train_sc)
15     print('depth = ',i,'Train Score',train_sc,'test Score',test_sc)
16 plt.plot(depths,train_scores,label='Train Score')
17 plt.plot(depths,test_scores,label='Test Score')
18 plt.xlabel('Depth')
19 plt.ylabel('Score')
20 plt.title('Depth vs score at depth of 5 at estimators = 115')
21 plt.show()

```

```

depth = 3 Train Score 0.8916120853581238 test Score 0.8687934859875491
depth = 9 Train Score 0.9572226298198419 test Score 0.9222953031452904
depth = 11 Train Score 0.9623451340902863 test Score 0.9252318758281279
depth = 15 Train Score 0.9634267621927706 test Score 0.9231288356496615
depth = 20 Train Score 0.9631629153051491 test Score 0.9235051024711141
depth = 35 Train Score 0.9634333127085721 test Score 0.9235601652753184
depth = 50 Train Score 0.9634333127085721 test Score 0.9235601652753184
depth = 70 Train Score 0.9634333127085721 test Score 0.9235601652753184
depth = 130 Train Score 0.9634333127085721 test Score 0.9235601652753184

```



```
In [0]: 1 from sklearn.metrics import f1_score
2 from sklearn.ensemble import RandomForestClassifier
3 from sklearn.metrics import f1_score
4 from sklearn.model_selection import RandomizedSearchCV
5 from scipy.stats import randint as sp_randint
6 from scipy.stats import uniform
7
8 param_dist = {"n_estimators": sp_randint(105, 125),
9              "max_depth": sp_randint(10, 15),
10             "min_samples_split": sp_randint(110, 190),
11             "min_samples_leaf": sp_randint(25, 65)}
12
13 clf = RandomForestClassifier(random_state=25, n_jobs=-1)
14
15 rf_random = RandomizedSearchCV(clf, param_distributions=param_dist,
16                               n_iter=5, cv=10, scoring='f1', random_state=25)
17
18 rf_random.fit(df_final_train, y_train)
19 print('mean test scores', rf_random.cv_results_['mean_test_score'])
20 print('mean train scores', rf_random.cv_results_['mean_train_score'])
```

```
mean test scores [0.96225043 0.96215493 0.96057081 0.96194015 0.96330005]
mean train scores [0.96294922 0.96266735 0.96115674 0.96263457 0.96430539]
```

```
In [0]: 1 print(rf_random.best_estimator_)
```

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',  
                        max_depth=14, max_features='auto', max_leaf_nodes=None,  
                        min_impurity_decrease=0.0, min_impurity_split=None,  
                        min_samples_leaf=28, min_samples_split=111,  
                        min_weight_fraction_leaf=0.0, n_estimators=121, n_jobs=-1,  
                        oob_score=False, random_state=25, verbose=0, warm_start=False)
```

```
In [0]: 1 clf = RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',  
2                                     max_depth=14, max_features='auto', max_leaf_nodes=None,  
3                                     min_impurity_decrease=0.0, min_impurity_split=None,  
4                                     min_samples_leaf=28, min_samples_split=111,  
5                                     min_weight_fraction_leaf=0.0, n_estimators=121, n_jobs=-1,  
6                                     oob_score=False, random_state=25, verbose=0, warm_start=False)
```

```
In [0]: 1 clf.fit(df_final_train,y_train)  
2 y_train_pred = clf.predict(df_final_train)  
3 y_test_pred = clf.predict(df_final_test)
```

```
In [0]: 1 from sklearn.metrics import f1_score  
2 print('Train f1 score',f1_score(y_train,y_train_pred))  
3 print('Test f1 score',f1_score(y_test,y_test_pred))
```

```
Train f1 score 0.9652533106548414
```


```
Test f1 score 0.9241678239279553
```

```
In [49]: 1 from sklearn.metrics import confusion_matrix
2 def plot_confusion_matrix(test_y, predict_y):
3     C = confusion_matrix(test_y, predict_y)
4
5     A = ((C.T)/(C.sum(axis=1))).T
6
7     B = (C/C.sum(axis=0))
8     plt.figure(figsize=(20,4))
9
10    labels = [0,1]
11    # representing A in heatmap format
12    cmap=sns.light_palette("blue")
13    plt.subplot(1, 3, 1)
14    sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
15    plt.xlabel('Predicted Class')
16    plt.ylabel('Original Class')
17    plt.title("Confusion matrix")
18
19    plt.subplot(1, 3, 2)
20    sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
21    plt.xlabel('Predicted Class')
22    plt.ylabel('Original Class')
23    plt.title("Precision matrix")
24
25    plt.subplot(1, 3, 3)
26    # representing B in heatmap format
27    sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
28    plt.xlabel('Predicted Class')
29    plt.ylabel('Original Class')
30    plt.title("Recall matrix")
31
32    plt.show()
```




```
In [0]: 1 print('Train confusion_matrix')
2 plot_confusion_matrix(y_train,y_train_pred)
3 print('Test confusion_matrix')
4 plot_confusion_matrix(y_test,y_test_pred)
```


Train confusion_matrix




Test confusion_matrix



```
In [0]: 1 from sklearn.metrics import roc_curve, auc
2 fpr,tpr,ths = roc_curve(y_test,y_test_pred)
3 auc_sc = auc(fpr, tpr)
4 plt.plot(fpr, tpr, color='navy',label='ROC curve (area = %0.2f)' % auc_sc)
5 plt.xlabel('False Positive Rate')
6 plt.ylabel('True Positive Rate')
7 plt.title('Receiver operating characteristic with test data')
8 plt.legend()
9 plt.show()
```



```
In [0]: 1 features = df_final_train.columns
        2 importances = clf.feature_importances_
        3 indices = (np.argsort(importances))[-25:]
        4 plt.figure(figsize=(10,12))
        5 plt.title('Feature Importances')
        6 plt.barh(range(len(indices)), importances[indices], color='r', align='center')
        7 plt.yticks(range(len(indices)), [features[i] for i in indices])
        8 plt.xlabel('Relative Importance')
        9 plt.show()
```




Assignments:

1. Add another feature called Preferential Attachment with followers and followees data of vertex. you can check about Preferential Attachment in below link <http://be.amazd.com/link-prediction/> (<http://be.amazd.com/link-prediction/>)
2. Add feature called svd_dot. you can calculate svd_dot as Dot product between source node svd and destination node svd features. you can read about this in below pdf https://storage.googleapis.com/kaggle-forum-message-attachments/2594/supervised_link_prediction.pdf (https://storage.googleapis.com/kaggle-forum-message-attachments/2594/supervised_link_prediction.pdf)
3. Tune hyperparameters for XG boost with all these features and check the error metric.

Preferential Attachment :

One well-known concept in social networks is that users with many friends tend to create more connections in the future. This is due to the fact that in some social networks, like in finance, the rich get richer. We estimate how "rich" our two vertices are by calculating the multiplication between the number of friends ($|\Gamma(x)|$) or followers each vertex has. It may be noted that the similarity index does not require any node neighbor information; therefore, this similarity index has the lowest computational complexity.



```
In [19]: 1 df_final_train.columns
        2 # df_final_train["num_followees_d"]
```

```
Out[19]: Index(['jaccard_followers', 'jaccard_followees', 'cosine_followers',
               'cosine_followees', 'num_followers_s', 'num_followees_s',
               'num_followees_d', 'inter_followers', 'inter_followees', 'adar_index',
               'follows_back', 'same_comp', 'shortest_path', 'weight_in', 'weight_out',
               'weight_f1', 'weight_f2', 'weight_f3', 'weight_f4', 'page_rank_s',
               'page_rank_d', 'katz_s', 'katz_d', 'hubs_s', 'hubs_d', 'authorities_s',
               'authorities_d', 'svd_u_s_1', 'svd_u_s_2', 'svd_u_s_3', 'svd_u_s_4',
               'svd_u_s_5', 'svd_u_s_6', 'svd_u_d_1', 'svd_u_d_2', 'svd_u_d_3',
               'svd_u_d_4', 'svd_u_d_5', 'svd_u_d_6', 'svd_v_s_1', 'svd_v_s_2',
               'svd_v_s_3', 'svd_v_s_4', 'svd_v_s_5', 'svd_v_s_6', 'svd_v_d_1',
               'svd_v_d_2', 'svd_v_d_3', 'svd_v_d_4', 'svd_v_d_5', 'svd_v_d_6',
               'num_followers_d'],
              dtype='object')
```

```
In [12]: 1 # Kind of like friends of friends and mutual friends
        2 # Perferntial Attachment is two vertices are by calculating the multiplication
        3 # between the number of friends ( $|\Gamma(x)|$ ) or followers each vertex has.
        4 # So Source and Destination Followees multiplication
```

```
In [20]: 1 df_final_train['perferntial_attach'] = df_final_train["num_followees_s"] * df_final_train["num_followees_d"]
2 df_final_train['perferntial_attach_follower'] = df_final_train["num_followers_s"] * df_final_train["num_fol
3
4
5 df_final_train.head()
```

Out[20]:

svd_v_s_6	svd_v_d_1	svd_v_d_2	svd_v_d_3	svd_v_d_4	svd_v_d_5	svd_v_d_6	num_followers_d	perferntial_attach	perferntial_attach
1.719702e-14	-1.355368e-12	4.675307e-13	1.128591e-06	6.616550e-14	9.771077e-13	4.159752e-14	6	120	
2.251737e-10	1.245101e-12	-1.636948e-10	-3.112650e-10	6.738902e-02	2.607801e-11	2.372904e-09	94	8662	
3.365389e-19	-1.238370e-18	1.438175e-19	-1.852863e-19	-5.901864e-19	1.629341e-19	-2.572452e-19	28	902	
4.498061e-13	-9.818087e-10	3.454672e-11	5.213635e-08	9.595823e-13	3.047045e-10	1.246592e-13	11	35	
1.407670e-14	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	1	33	

```
In [21]: 1 df_final_test['perferntial_attach'] = df_final_test["num_followees_s"] * df_final_test["num_followees_d"]
2 df_final_test['perferntial_attach_follower'] = df_final_test["num_followers_s"] * df_final_test["num_followers_d"]
3 df_final_test.head()
```

Out[21]:

inter_followees	adar_index	...	svd_v_s_6	svd_v_d_1	svd_v_d_2	svd_v_d_3	svd_v_d_4	svd_v_d_5	svd_v_d_6	num_followers_d	perferntial_attach
0	0.000000	...	5.535503e-14	-9.994076e-10	5.791910e-10	3.512364e-07	2.486658e-09	2.771146e-09	1.727694e-12	14	54
0	0.000000	...	4.701436e-15	-9.360516e-12	3.206809e-10	4.668696e-08	6.665777e-12	1.495979e-10	9.836670e-14	17	19
0	0.000000	...	4.199834e-14	-4.253075e-13	4.789463e-13	3.479824e-07	1.630549e-13	3.954708e-13	3.875785e-14	10	144
0	0.000000	...	2.817657e-13	-2.162590e-11	6.939194e-12	1.879861e-05	4.384816e-12	1.239414e-11	6.483485e-13	37	340
7	6.136433	...	9.656662e-14	-8.742904e-12	7.467370e-12	1.256880e-05	3.636983e-12	3.948463e-12	2.415863e-13	27	405

Adding SVD Dot Feature that is product of source and destination SVDs

```
In [ ]: 1 # Dot product of Source and Destination Node SVD
```

```
In [25]: 1 source = ['svd_u_s_1', 'svd_u_s_2', 'svd_u_s_3', 'svd_u_s_4', 'svd_u_s_5', 'svd_u_s_6']
2 source_v = ['svd_v_s_1', 'svd_v_s_2', 'svd_v_s_3', 'svd_v_s_4', 'svd_v_s_5', 'svd_v_s_6']
3 destination = ['svd_u_d_1', 'svd_u_d_2', 'svd_u_d_3', 'svd_u_d_4', 'svd_u_d_5', 'svd_u_d_6']
4 destination_v = ['svd_v_d_1', 'svd_v_d_2', 'svd_v_d_3', 'svd_v_d_4', 'svd_v_d_5', 'svd_v_d_6']
5
```



```
In [22]: 1 # source_array = []
2 # destination_array = []
3 # svd_dot = []
4
5 # for num in tqdm(range(df_final_train.shape[0])):
6 #     for i,j in (zip(source,destination)):
7 #         source_array.append(np.array(df_final_train[i].iloc[num]))
8 #         destination_array.append(np.array(df_final_train[j].iloc[num]))
9 #         svd_dot.append(np.dot(source_array, destination_array))
10 # df_final_train['svd_dot']=svd_dot
11
```

```
In [32]: 1 # Corrections
2 f = 0;
3 for i,j in tqdm(zip(source,destination)):
4     f = f + df_final_train[i] * df_final_train[j]
5 df_final_train['svd_dot_source'] = f
6
7 f = 0;
8 for i,j in tqdm(zip(source_v,destination_v)):
9     f = f + df_final_train[i] * df_final_train[j]
10 df_final_train['svd_dot_dest'] = f
11
12 f = 0;
13 for i,j in tqdm(zip(source,destination)):
14     f = f + df_final_test[i] * df_final_test[j]
15 df_final_test['svd_dot_source'] = f
16
17 f = 0;
18 for i,j in tqdm(zip(source_v,destination_v)):
19     f = f + df_final_test[i] * df_final_test[j]
20 df_final_test['svd_dot_dest'] = f
21
```

6it [00:00, 482.16it/s]

6it [00:00, 480.27it/s]

6it [00:00, 623.44it/s]

6it [00:00, 337.19it/s]

```
In [33]: 1 df_final_train.head()
```

```
Out[33]:
```

id_v_d_2	svd_v_d_3	svd_v_d_4	svd_v_d_5	svd_v_d_6	num_followers_d	perferntial_attach	perferntial_attach_follower	svd_dot_source	svd
3307e-13	1.128591e-06	6.616550e-14	9.771077e-13	4.159752e-14	6	120	36	1.114958e-11	2.23
336948e-10	-3.112650e-10	6.738902e-02	2.607801e-11	2.372904e-09	94	8662	8836	3.192812e-03	9.06
3175e-19	-1.852863e-19	-5.901864e-19	1.629341e-19	-2.572452e-19	28	902	784	1.787503e-35	2.46
1672e-11	5.213635e-08	9.595823e-13	3.047045e-10	1.246592e-13	11	35	121	4.710376e-20	3.11
000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	1	33	1	7.773952e-14	0.00

```
In [35]: 1 df_final_test.head()
```

```
Out[35]:
```

	jaccard_followers	jaccard_followees	cosine_followers	cosine_followees	num_followers_s	num_followees_s	num_followees_d	inter_followers	inter
0	0	0.0	0.029161	0.000000	14	6	9	1	
1	0	0.0	0.000000	0.000000	17	1	19	0	
2	0	0.0	0.000000	0.000000	10	16	9	0	
3	0	0.0	0.000000	0.000000	37	10	34	0	
4	0	0.2	0.042767	0.347833	27	15	27	4	

5 rows × 56 columns

```
In [ ]: 1
```

Random Forest


```

In [37]: 1 from sklearn.metrics import f1_score
2 from sklearn.ensemble import RandomForestClassifier
3 from sklearn.metrics import f1_score
4 from sklearn.model_selection import RandomizedSearchCV
5 from scipy.stats import randint as sp_randint
6 from scipy.stats import uniform
7
8 param_dist = {"n_estimators": sp_randint(105, 125),
9              "max_depth": sp_randint(10, 15)}
10
11
12 clf = RandomForestClassifier(random_state=25, n_jobs=-1)
13
14 rf_random = RandomizedSearchCV(clf, param_distributions=param_dist,
15                               cv=2, scoring='f1')
16
17 rf_random.fit(df_final_train, y_train)
18

```

```

Out[37]: RandomizedSearchCV(cv=2, error_score='raise-deprecating',
                             estimator=RandomForestClassifier(bootstrap=True,
                                                                class_weight=None,
                                                                criterion='gini',
                                                                max_depth=None,
                                                                max_features='auto',
                                                                max_leaf_nodes=None,
                                                                min_impurity_decrease=0.0,
                                                                min_impurity_split=None,
                                                                min_samples_leaf=1,
                                                                min_samples_split=2,
                                                                min_weight_fraction_leaf=0.0,
                                                                n_estimators='warn',
                                                                n_jobs=-1, oob_score...
                                                                random_state=25, verbose=0,
                                                                warm_start=False),
                             iid='warn', n_iter=10, n_jobs=None,
                             param_distributions={'max_depth': <scipy.stats._distn_infrastructure.rv_frozen object at 0
x1a3f8af7f0>,
                                                  'n_estimators': <scipy.stats._distn_infrastructure.rv_frozen object a
t 0x1a3f9455c0>},

```

```
pre_dispatch='2*n_jobs', random_state=None, refit=True,  
return_train_score=False, scoring='f1', verbose=0)
```

```
In [38]: 1 print('mean test scores',rf_random.cv_results_['mean_test_score'])
```

```
mean test scores [0.9663106  0.96309182 0.96826787 0.9609609  0.96657034 0.96474245  
0.96476606 0.96826012 0.96827655 0.96651507]
```

```
In [39]: 1 print(rf_random.best_estimator_)
```

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',  
                        max_depth=14, max_features='auto', max_leaf_nodes=None,  
                        min_impurity_decrease=0.0, min_impurity_split=None,  
                        min_samples_leaf=1, min_samples_split=2,  
                        min_weight_fraction_leaf=0.0, n_estimators=114,  
                        n_jobs=-1, oob_score=False, random_state=25, verbose=0,  
                        warm_start=False)
```

```
In [44]: 1 n_estimator_rf = 114  
2 max_depth_rf = 14  
3 clf = RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',  
4                             max_depth=14, max_features='auto', max_leaf_nodes=None,  
5                             min_impurity_decrease=0.0, min_impurity_split=None,  
6                             min_samples_leaf=1, min_samples_split=2,  
7                             min_weight_fraction_leaf=0.0, n_estimators=114,  
8                             n_jobs=-1, oob_score=False, random_state=25, verbose=0,  
9                             warm_start=False)
```

```
In [45]: 1 clf.fit(df_final_train,y_train)  
2 y_train_pred = clf.predict(df_final_train)  
3 y_test_pred = clf.predict(df_final_test)
```


```
In [46]: 1 from sklearn.metrics import f1_score
          2
          3 Train_f1_Score_rf = f1_score(y_train,y_train_pred)
          4 Test_f1_Score_rf = f1_score(y_test,y_test_pred)
          5
          6 print('Train f1 score',Train_f1_Score_rf)
          7 print('Test f1 score',Test_f1_Score_rf)
```

Train f1 score 0.9720793459917744


Test f1 score 0.9188237291012475

```
In [50]: 1 print('Train confusion_matrix')
2 plot_confusion_matrix(y_train,y_train_pred)
3 print('Test confusion_matrix')
4 plot_confusion_matrix(y_test,y_test_pred)
```

Train confusion_matrix



Test confusion_matrix



```
In [ ]: 1 features = df_final_train.columns
2 importances = clf.feature_importances_
3 indices = (np.argsort(importances))[-25:]
4 plt.figure(figsize=(10,12))
5 plt.title('Feature Importances')
6 plt.barh(range(len(indices)), importances[indices], color='r', align='center')
7 plt.yticks(range(len(indices)), [features[i] for i in indices])
8 plt.xlabel('Relative Importance')
9 plt.show()
```

XGboost

```
In [40]: 1 from datetime import datetime
2 from sklearn.metrics import f1_score
3 from sklearn.model_selection import RandomizedSearchCV
4 from xgboost import XGBClassifier
```

```
In [42]: 1 start = datetime.now()
2 clf = XGBClassifier()
3 params = {"n_estimators":sp_randint(105,125),
4           "max_depth": sp_randint(10,15)
5           }
6 rf = RandomizedSearchCV(clf, param_distributions=params,cv=2,scoring='f1', n_jobs=-1)
7
8
9
10 rf.fit(df_final_train,y_train)
11 print('mean test scores',rf.cv_results_['mean_test_score'])
12 print(datetime.now() - start)
```

```
mean test scores [0.97926699 0.97903336 0.97881123 0.97901685 0.9791557  0.97929291
0.97896439 0.97933791 0.97906177 0.97911945]
0:14:49.802030
```



```
In [43]: 1 print(rf.best_estimator_)
```

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
              colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,  
              max_depth=12, min_child_weight=1, missing=None, n_estimators=118,  
              n_jobs=1, nthread=None, objective='binary:logistic',  
              random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,  
              seed=None, silent=True, subsample=1)
```

```




In [51]: 1 n_estimator_xgb = 118
2 max_depth_xgb = 12
3 xgb = XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
4                     colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
5                     max_depth=12, min_child_weight=1, missing=None, n_estimators=118,
6                     n_jobs=1, nthread=None, objective='binary:logistic',
7                     random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
8                     seed=None, silent=True, subsample=1)
9 xgb.fit(df_final_train,y_train)
10 y_pred_train = xgb.predict(df_final_train)
11 y_pred_test = xgb.predict(df_final_test)
12
13 Train_f1_Score_xgb = f1_score(y_train,y_pred_train)
14 Test_f1_Score_xgb = f1_score(y_test,y_pred_test)
15
16 print('Train f1 score',Train_f1_Score_xgb)
17 print('Test f1 score',Test_f1_Score_xgb)
18
19 print('Train confusion_matrix')
20 plot_confusion_matrix(y_train,y_pred_train)
21 print('Test confusion_matrix')
22 plot_confusion_matrix(y_test,y_pred_test)

```


Train f1 score 0.9981790895447723

Test f1 score 0.9263278611931106


Train confusion_matrix




Test confusion_matrix



```
In [52]: 1 from sklearn.metrics import roc_curve, auc
2 fpr,tpr,ths = roc_curve(y_test,y_pred_test)
3 auc_sc = auc(fpr, tpr)
4 plt.plot(fpr, tpr, color='navy',label='ROC curve (area = %0.2f)' % auc_sc)
5 plt.xlabel('False Positive Rate')
6 plt.ylabel('True Positive Rate')
7 plt.title('Receiver operating characteristic with test data')
8 plt.legend()
9 plt.show()
```



```
In [53]: 1 features = df_final_train.columns
          2 importances = xgb.feature_importances_
          3 indices = (np.argsort(importances))[-25:]
          4 plt.figure(figsize=(10,12))
          5 plt.title('Feature Importances')
          6 plt.barh(range(len(indices)), importances[indices], color='r', align='center')
          7 plt.yticks(range(len(indices)), [features[i] for i in indices])
          8 plt.xlabel('Relative Importance')
          9 plt.show()
```



```
In [54]: 1 from prettytable import PrettyTable
2 x = PrettyTable()
3 x.field_names = ["Model", "n_estimators", "max_depth", "Train_f1_Score", "Test_f1_Score"]
4 x.add_row(['Random Forest', n_estimator_rf, max_depth_rf, Train_f1_Score_rf, Test_f1_Score_rf])
5 x.add_row(['XGBOOST', n_estimator_xgb, max_depth_xgb, Train_f1_Score_xgb, Test_f1_Score_xgb])
6 print(x)
```

Model	n_estimators	max_depth	Train_f1_Score	Test_f1_Score
Random Forest	114	14	0.9720793459917744	0.9188237291012475
XGBOOST	118	12	0.9981790895447723	0.9263278611931106

Conclusion :

- Facebook Friend Recommendation Case study is one of its kind as most of the Internet Companies like Facebook, Instagram, reddit, Github etc has graph based features. Understanding various Graph feature is important and all through a fun ride.
- In EDA we made use of np.percentile and saw how we can even look at 99.1 percentile
- Explored different Graph Based Feature Engineering in which **Follow back & cosine followers** were important one
- Added new features as part of Assignment
- XGboost performed better than Random Forest

```
In [ ]: 1
```