

**Try out models (Logistic regression, Linear-SVM) with simple TF-IDF vectors instead of TD\_IDF weighted word2Vec.**

```
In [1]: 1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import re
4 import time
5 import warnings
6 import sqlite3
7 from sqlalchemy import create_engine # database connection
8 import csv
9 import os
10 warnings.filterwarnings("ignore")
11 import datetime as dt
12 import numpy as np
13 from nltk.corpus import stopwords
14 from sklearn.decomposition import TruncatedSVD
15 from sklearn.preprocessing import normalize
16 from sklearn.feature_extraction.text import CountVectorizer
17 from sklearn.manifold import TSNE
18 import seaborn as sns
19 from sklearn.neighbors import KNeighborsClassifier
20 from sklearn.metrics import confusion_matrix
21 from sklearn.metrics.classification import accuracy_score, log_loss
22 from sklearn.feature_extraction.text import TfidfVectorizer
23 from collections import Counter
24 from scipy.sparse import hstack
25 from sklearn.multiclass import OneVsRestClassifier
26 from sklearn.svm import SVC
27 from sklearn.model_selection import StratifiedKFold
28 from collections import Counter, defaultdict
29 from sklearn.calibration import CalibratedClassifierCV
30 from sklearn.naive_bayes import MultinomialNB
31 from sklearn.naive_bayes import GaussianNB
32 from sklearn.model_selection import train_test_split
33 from sklearn.model_selection import GridSearchCV
34 import math
35 from sklearn.metrics import normalized_mutual_info_score
36 from sklearn.ensemble import RandomForestClassifier
37
38
39
40 from sklearn.model_selection import cross_val_score
41 from sklearn.linear_model import SGDClassifier
```

```

42 from mlxtend.classifier import StackingClassifier
43
44 from sklearn import model_selection
45 from sklearn.linear_model import LogisticRegression
46 from sklearn.metrics import precision_recall_curve, auc, roc_curve

```

```
In [163]: 1 rows = 400000
```

```

In [164]: 1 if os.path.isfile('nlp_features_train.csv'):
2         dfnlp = pd.read_csv("nlp_features_train.csv",encoding='latin-1', nrows = rows)
3 else:
4     print("download nlp_features_train.csv from drive or run previous notebook")
5
6 if os.path.isfile('df_fe_without_preprocessing_train.csv'):
7     dfppro = pd.read_csv("df_fe_without_preprocessing_train.csv",encoding='latin-1', nrows = rows)
8 else:
9     print("download df_fe_without_preprocessing_train.csv from drive or run previous notebook")

```

```
In [165]: 1 dfnlp.head(2)
```

Out[165]:

	id	qid1	qid2	question1	question2	is_duplicate	cwc_min	cwc_max	csc_min	csc_max	...	ctc_max	last_word_eq	first_word_eq	abs_len_di
0	0	1	2	what is the step by step guide to invest in sh...	what is the step by step guide to invest in sh...	0	0.999980	0.833319	0.999983	0.999983	...	0.785709	0.0	1.0	2.
1	1	3	4	what is the story of kohinoor koh i noor dia...	what would happen if the indian government sto...	0	0.799984	0.399996	0.749981	0.599988	...	0.466664	0.0	1.0	5.

2 rows × 21 columns

```
In [166]: 1 print(dfnlp.shape)
          2 print(dfppro.shape)
```

```
(400000, 21)
(400000, 17)
```

```
In [167]: 1 dfppro.head(2)
```

Out[167]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total	wor
0	0	1	2	What is the step by step guide to invest in sh...	What is the step by step guide to invest in sh...	0	1	1	66	57	14	12	10.0	23.0	0
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	What would happen if the Indian government sto...	0	4	1	51	88	8	13	4.0	20.0	0

**Have to combine these 2 dfs first So that later we can add just one tfidf vector for questions via hstack**

**As question1 and ques2 are preprocessed in dfnlp then we must vectorize with dfnlp questions**

```
In [168]: 1 dfppro = dfppro.drop(['qid1', 'qid2', 'question1', 'question2', 'is_duplicate'], axis=1)
          2 dfppro.head()
```

Out[168]:

	id	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total	word_share	freq_q1+q2	freq_q1-q2
0	0	1	1	66	57	14	12	10.0	23.0	0.434783	2	0
1	1	4	1	51	88	8	13	4.0	20.0	0.200000	5	3
2	2	1	1	73	59	14	10	4.0	24.0	0.166667	2	0
3	3	1	1	50	65	11	9	0.0	19.0	0.000000	2	0
4	4	3	1	76	39	13	7	2.0	20.0	0.100000	4	2

```
In [169]: 1 df = dfnlp.merge(dfppro, on=['id'], how='left')
```

```
In [ ]: 1
```

```
In [170]: 1 print(df.shape)

(400000, 32)
```

**Now we have everything preprocessed, So lets first split our data before doing vectorizing as it might cause data leakage which we dont want at any cost**

```
In [171]: 1 df.columns
```

```
Out[171]: Index(['id', 'qid1', 'qid2', 'question1', 'question2', 'is_duplicate',
               'cwc_min', 'cwc_max', 'csc_min', 'csc_max', 'ctc_min', 'ctc_max',
               'last_word_eq', 'first_word_eq', 'abs_len_diff', 'mean_len',
               'token_set_ratio', 'token_sort_ratio', 'fuzz_ratio',
               'fuzz_partial_ratio', 'longest_substr_ratio', 'freq_qid1', 'freq_qid2',
               'q1len', 'q2len', 'q1_n_words', 'q2_n_words', 'word_Common',
               'word_Total', 'word_share', 'freq_q1+q2', 'freq_q1-q2'],
              dtype='object')
```

```
In [ ]: 1
```

```
In [172]: 1 df["question1"].head()
```

```
Out[172]: 0    what is the step by step guide to invest in sh...
          1    what is the story of kohinoor koh i noor dia...
          2    how can i increase the speed of my internet co...
          3    why am i mentally very lonely how can i solve...
          4    which one dissolve in water quikly sugar salt...
          Name: question1, dtype: object
```

```
In [173]: 1 y_true = df["is_duplicate"]
```

```
In [174]: 1 df.drop(['is_duplicate', 'qid1', 'qid2'], axis=1, inplace=True)
```

```
In [175]: 1 print(df.shape)
```

```
(400000, 29)
```

```
In [176]: 1 df = df.fillna(' ')
```

```
In [ ]: 1
```

```
In [177]: 1 X_train,X_test, y_train, y_test = train_test_split(df, y_true, stratify=y_true, test_size=0.3)
          2 print("Number of data points in train data :",X_train.shape)
          3 print("Number of data points in test data :",X_test.shape)
```

```
Number of data points in train data : (280000, 29)
```

```
Number of data points in test data : (120000, 29)
```

```
In [178]: 1 from sklearn.feature_extraction.text import TfidfVectorizer
2 questions_train = list(X_train['question1']) + list(X_train['question2'])
3
4 tfidf = TfidfVectorizer(lowercase=False)
5 tfidf.fit(questions_train)
6
7 tfidf_train_q1 = tfidf.transform(X_train['question1'])
8 tfidf_train_q2 = tfidf.transform(X_train['question2'])
9 tfidf_test_q1 = tfidf.transform(X_test['question1'])
10 tfidf_test_q2 = tfidf.transform(X_test['question2'])
11
12 print(tfidf_train_q1.shape)
13 print(tfidf_train_q2.shape)
14 print(tfidf_test_q1.shape)
15 print(tfidf_test_q2.shape)
16
17
```

```
(280000, 73371)
(280000, 73371)
(120000, 73371)
(120000, 73371)
```

```
In [179]: 1 X_train = X_train.drop(['question1', 'question2'], axis=1)
2 X_test = X_test.drop(['question1', 'question2'], axis=1)
```

```
In [180]: 1 print(X_train.shape)
```

```
(280000, 27)
```

```
In [181]: 1 from scipy.sparse import hstack
2 X_train_final = hstack([tfidf_train_q1, tfidf_train_q2, X_train.values]).tocsr()
3 X_test_final = hstack([tfidf_test_q1, tfidf_test_q2, X_test.values]).tocsr()
4 print(X_train_final.shape, X_test_final.shape)
```

```
(280000, 146769) (120000, 146769)
```

```

In [182]: 1 # This function plots the confusion matrices given y_i, y_i_hat.
2 def plot_confusion_matrix(test_y, predict_y):
3     C = confusion_matrix(test_y, predict_y)
4     # C = 9,9 matrix, each cell (i,j) represents number of points of class i are predicted class j
5
6     A = ((C.T)/(C.sum(axis=1))).T
7     #divid each element of the confusion matrix with the sum of elements in that column
8
9     # C = [[1, 2],
10    #      [3, 4]]
11    # C.T = [[1, 3],
12    #        [2, 4]]
13    # C.sum(axis = 1)  axis=0 corresponds to columns and axis=1 corresponds to rows in two dimensional array
14    # C.sum(axix =1) = [[3, 7]]
15    # ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7]
16    #                             [2/3, 4/7]]
17
18    # ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3]
19    #                             [3/7, 4/7]]
20    # sum of row elements = 1
21
22    B = (C/C.sum(axis=0))
23    #divid each element of the confusion matrix with the sum of elements in that row
24    # C = [[1, 2],
25    #      [3, 4]]
26    # C.sum(axis = 0)  axis=0 corresponds to columns and axis=1 corresponds to rows in two dimensional array
27    # C.sum(axix =0) = [[4, 6]]
28    # (C/C.sum(axis=0)) = [[1/4, 2/6],
29    #                       [3/4, 4/6]]
30    plt.figure(figsize=(20,4))
31
32    labels = [0,1]
33    # representing A in heatmap format
34    cmap=sns.light_palette("blue")
35    plt.subplot(1, 3, 1)
36    sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
37    plt.xlabel('Predicted Class')
38    plt.ylabel('Original Class')
39    plt.title("Confusion matrix")
40
41    plt.subplot(1, 3, 2)

```



```

42     sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
43     plt.xlabel('Predicted Class')
44     plt.ylabel('Original Class')
45     plt.title("Precision matrix")
46
47     plt.subplot(1, 3, 3)
48     # representing B in heatmap format
49     sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
50     plt.xlabel('Predicted Class')
51     plt.ylabel('Original Class')
52     plt.title("Recall matrix")
53
54     plt.show()

```

## Random Model

```

In [183]: 1 print("-"*10, "Distribution of output variable in train data", "-"*10)
2 train_distr = Counter(y_train)
3 train_len = len(y_train)
4 print("Class 0: ",int(train_distr[0])/train_len,"Class 1: ", int(train_distr[1])/train_len)
5 print("-"*10, "Distribution of output variable in train data", "-"*10)
6 test_distr = Counter(y_test)
7 test_len = len(y_test)
8 print("Class 0: ",int(test_distr[1])/test_len, "Class 1: ",int(test_distr[1])/test_len)

```


```

----- Distribution of output variable in train data -----
Class 0:  0.6305785714285714 Class 1:  0.36942142857142857
----- Distribution of output variable in train data -----
Class 0:  0.369425 Class 1:  0.369425

```

```
In [184]: 1 # we need to generate 9 numbers and the sum of numbers should be 1
2 # one solution is to generate 9 numbers and divide each of the numbers by their sum
3 # ref: https://stackoverflow.com/a/18662466/4084039
4 # we create a output array that has exactly same size as the CV data
5 predicted_y = np.zeros((test_len,2))
6 for i in range(test_len):
7     rand_probs = np.random.rand(1,2)
8     predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
9 print("Log loss on Test Data using Random Model",log_loss(y_test, predicted_y, eps=1e-15))
10
11 predicted_y = np.argmax(predicted_y, axis=1)
12 plot_confusion_matrix(y_test, predicted_y)
```

Log loss on Test Data using Random Model 0.8937041392914908



## LOGISTIC REGRESSION

```

In [192]: alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.
2
log_error_array=[]
for i in alpha:
5 clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
6 clf.fit(X_train_final, y_train)
7 sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
8 sig_clf.fit(X_train_final, y_train)
9 predict_y = sig_clf.predict_proba(X_test_final)
10 log_error_array.append(log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
11 print('For values of alpha = ', i, "The log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
12
fig, ax = plt.subplots()
ax.plot(alpha, log_error_array, c='g')
for i, txt in enumerate(np.round(log_error_array, 3)):
16 ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
22
23
best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(X_train_final, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train_final, y_train)
29
predict_y = sig_clf.predict_proba(X_train_final)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(X_test_final)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
predicted_y = np.argmax(predict_y, axis=1)
print("Total number of data points :", len(predicted_y))
plt.plot_confusion_matrix(y_test, predicted_y)

```

```


For values of alpha = 1e-05 The log loss is: 0.6586489466732767
For values of alpha = 0.0001 The log loss is: 0.6586489466732767
For values of alpha = 0.001 The log loss is: 0.6586489466732767
For values of alpha = 0.01 The log loss is: 0.614065123886705

```

For values of alpha = 0.1 The log loss is: 0.6144293652923921

For values of alpha = 1 The log loss is: 0.6586489466732767


For values of alpha = 10 The log loss is: 0.6574867280589907



For values of best alpha = 0.01 The train log loss is: 0.6143797442084187

For values of best alpha = 0.01 The test log loss is: 0.614065123886705

Total number of data points : 120000



```
In [193]: 1 LogisticRegression_alpha = alpha[best_alpha]
          2 LogisticRegression_logloss = log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15)
```

# LINEAR SVM

```

In [194]: 1 alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.
2
3 # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.
4 # -----
5 # default parameters
6 # SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None,
7 # shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, pow
8 # class_weight=None, warm_start=False, average=False, n_iter=None)
9
10 # some of methods
11 # fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
12 # predict(X) Predict class labels for samples in X.
13
14 #-----
15 # video link:
16 #-----
17
18
19 log_error_array=[]
20 for i in alpha:
21     clf = SGDClassifier(alpha=i, penalty='l1', loss='hinge', random_state=42)
22     clf.fit(X_train_final, y_train)
23     sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
24     sig_clf.fit(X_train_final, y_train)
25     predict_y = sig_clf.predict_proba(X_test_final)
26     log_error_array.append(log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
27     print('For values of alpha = ', i, "The log loss is:", log_loss(y_test, predict_y, labels=clf.classes_,
28
29 fig, ax = plt.subplots()
30 ax.plot(alpha, log_error_array, c='g')
31 for i, txt in enumerate(np.round(log_error_array, 3)):
32     ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], log_error_array[i]))
33 plt.grid()
34 plt.title("Cross Validation Error for each alpha")
35 plt.xlabel("Alpha i's")
36 plt.ylabel("Error measure")
37 plt.show()
38
39
40 best_alpha = np.argmin(log_error_array)
41 clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l1', loss='hinge', random_state=42)


```

```


42 clf.fit(X_train_final, y_train)
43 sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
44 sig_clf.fit(X_train_final, y_train)
45
46 predict_y = sig_clf.predict_proba(X_train_final)
47 print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y))
48 predict_y = sig_clf.predict_proba(X_test_final)
49 print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y))
50 predicted_y = np.argmax(predict_y, axis=1)
51 print("Total number of data points :", len(predicted_y))
52 plot_confusion_matrix(y_test, predicted_y)

```

For values of alpha = 1e-05 The log loss is: 0.6586489466732767  
 For values of alpha = 0.0001 The log loss is: 0.6586489466732767  
 For values of alpha = 0.001 The log loss is: 0.6586489466732767  
 For values of alpha = 0.01 The log loss is: 0.6586489466732767  
 For values of alpha = 0.1 The log loss is: 0.61342744507318  
 For values of alpha = 1 The log loss is: 0.5886888763223536  
 For values of alpha = 10 The log loss is: 0.6474710342982285



For values of best alpha = 1 The train log loss is: 0.5891724271157317  
 For values of best alpha = 1 The test log loss is: 0.5886888763223536  
 Total number of data points : 120000



```
In [195]: 1 SVM_alpha = alpha[best_alpha]
          2 SVM_logloss = log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15)
```

## SUMMARY

```
In [196]: 1 from prettytable import PrettyTable
          2
          3 x = PrettyTable()
          4 x.field_names = ["Vectorizer", "Model", "Hyperparameter (ALPHA)", "LOG LOSS", "Data Points"]
          5 x.add_row(["Tfidf", "Logistic Regression", LogisticRegression_alpha, LogisticRegression_logloss, rows])
          6 x.add_row(["Tfidf", "Linear SVM", SVM_alpha, SVM_logloss, rows])
          7 print(x)
```

Vectorizer	Model	Hyperparameter (ALPHA)	LOG LOSS	Data Points
Tfidf	Logistic Regression	0.01	0.614065123886705	400000
Tfidf	Linear SVM	1	0.5886888763223536	400000

**TFIDF weighted W2V (100k datapoints) gave better performance than TFIDF Vectorizer (400k datapoints).**



